DESIGN PROTOTYPE — SAMPLE EDITION. The organisations, people, quotes and figures in this edition are fictional and are not current news.

FRIDAY, 2 OCTOBER 2026 YOUR 4-MINUTE EDITION

Over Coffee

AI

Good morning. Here's what matters.

FIVE THINGS WORTH KNOWING TODAY

KESTREL 4 API PRICE	AI CHIP LEAD TIME	FIRMS RUNNING AI IN PRODUCTION
$1.20 per M tokens	9 weeks	38%
▼ from $3.00 · Kestrel 3	▼ from 26 · A year ago	▲ from 22% · A year ago

AI · MODEL RELEASES ● CONFIRMED

Kestrel Labs cuts prices 60% with Kestrel 4 — and gives a smaller model away

The new model family is cheaper to run and, on the company's own tests, better at code. The price, not the benchmark, is the story.



Kestrel Labs released Kestrel 4 on Thursday, pricing its flagship API at $1.20 per million input tokens, 60% below Kestrel 3. It also published Kestrel 4 Mini with open weights under a licence that permits commercial use.

The company says Kestrel 4 completes 71% of tasks on its internal coding benchmark, up from 54%. That figure has not been independently tested. The technical report includes outside safety evaluations but does not disclose training compute.

Hands-on testing by Circuit Review found the Mini model runs on a single 48GB workstation GPU, putting a capable model within reach of teams that cannot or will not send data to a cloud provider.

API PRICE PER MILLION INPUT TOKENS



Source: Kestrel Labs

TODAY IN BRIEF

1. Kestrel Labs cut its API price 60% with Kestrel 4 and released a smaller open-weights model.
2. A draft code would require independent audits before the most capable AI models are released.
3. An AI weather model beat physics-based forecasts at 10 days in an early, unreviewed study.
4. Varda says waits for AI chips are down to nine weeks from 26 — the shortage is easing.
5. AI in production rose to 38% of surveyed firms, but only 9% report savings.

ONE THING TO REMEMBER

Capable AI is getting cheaper faster than it is getting smarter. Budgets, not breakthroughs, will set the pace of adoption.

IN THIS EDITION

WHY IT MATTERS

For most companies, AI spending is limited by cost per task rather than by raw capability. A cut of this size changes which projects survive a budget review.

WHAT CHANGED

A frontier lab releasing open weights alongside a 60% price cut resets the floor for what buyers expect to pay.

WHAT HAPPENS NEXT

Analysts expect rivals to respond on price within weeks. Independent benchmark results should follow.

LEARN MORE overcoffee.app/s/preview-single-10/st-ai-kestrel-4

AI 5 stories

AI

LEAD STORY ON THE FRONT PAGE · 4 MORE STORIES

REGULATION ● CONFIRMED

Draft rules would require independent audits before major AI model releases

The Atlantic Digital Standards Board has published its first detailed code. The open question is who it covers.

The Atlantic Digital Standards Board published a draft code on Thursday that would require developers of the most capable general-purpose models to commission an independent audit before release and to publish a summary of the results.

Consultation closes on 15 December. The Board intends final rules to apply from 1 July 2027. Industry groups say the threshold that decides which models are covered is not yet clear enough to apply.

Capacity is a practical constraint: the Auditing Standards Forum says fewer than 10 organisations can currently conduct audits of this kind.

WHY IT MATTERS

Pre-release audits would add time and cost to every major launch, and create a market for auditors that barely exists today.

WHAT HAPPENS NEXT

Expect detailed industry submissions before the December deadline, most of them focused on the threshold.

LEARN MORE overcoffee.app/s/preview-single-10/st-ai-audit-code

RESEARCH · FORECASTING

○ EMERGING · NOT YET PEER REVIEWED

A learned weather model beats physics on 10-day forecasts, in a study not yet peer reviewed

The Northern Meteorological Service will run it alongside its existing model from November. Growers are watching closely.

A machine-learning weather model developed by the Northern Meteorological Service and Aster Systems was more accurate than the leading physics-based model on 87% of the variables tested over a two-year evaluation, according to a preprint published this week. The paper has not yet been peer reviewed.

The service will run the model operationally alongside its physics model from November, while continuing to issue official forecasts from the physics model.

For agriculture, the useful horizon matters most. Agronomists say a reliable 10-day window would change how growers time planting, spraying and harvest.

WHY IT MATTERS

Medium-range forecasts drive costly decisions in farming, energy and logistics. A real improvement in the 10-day window has value well beyond meteorology.

WHAT HAPPENS NEXT

Peer review, and the service's own side-by-side evaluation over the winter.

LEARN MORE overcoffee.app/s/preview-single-10/st-ai-weather-model

HARDWARE · CONFIRMED

Varda says AI chip lead times have fallen to nine weeks, from 26

The shortage that shaped two years of AI spending is easing. Smaller buyers stand to gain most.

Varda Semiconductor said lead times for its V-series AI accelerators have fallen to nine weeks, from 26 weeks a year ago. Data-centre revenue rose 41% year on year to $8.4bn.

The company expects supply to stay balanced through the first half of 2027. Shares rose 6% in after-hours trading, Meridian Wire reported.

LEAD TIME FOR V-SERIES ACCELERATORS, WEEKS



Source: Varda Semiconductor

WHY IT MATTERS

Access to chips, more than money or talent, has decided who could build and run large AI systems. Shorter queues open the door to smaller cloud providers and enterprises that could not secure allocation.

LEARN MORE overcoffee.app/s/preview-single-10/st-ai-chip-lead-times

ENTERPRISE ADOPTION · CONFIRMED

38% of companies now run AI in production. Only 9% can show savings

Adoption is accelerating. Proof of return is not.

A Halberd Institute survey of 1,200 companies found 38% have at least one AI system in production, up from 22% a year earlier. Only 9% report measurable cost savings at company level.

Data quality was the most cited barrier (61%), ahead of cost (47%). The results are self-reported.

COMPANIES WITH AI IN PRODUCTION, %



Source: Halberd Institute

WHY IT MATTERS

The gap between deployment and measurable savings is where AI budgets will be tested in next year's planning rounds.

LEARN MORE overcoffee.app/s/preview-single-10/st-ai-enterprise-survey

That's today's edition. You're done.

Prepared 10:26 UTC · 23 sources from 16 publishers reviewed · 8 developments considered · 5 published · Quality-checked, 11 of 11 checks passed